

CNN and Vision Transformer Pre-trained Models for Facial Emotion Recognition

Group 2

Luoning Zhang, Tom Gao, William Chu, Montek Kalsi, and Grant Liu



Introduction


Our project explores two distinct architectures, CNNs and Vision Transformers, to recognize facial emotions.

Why emotion recognition? It's a useful and relevant problem for many modern day purposes

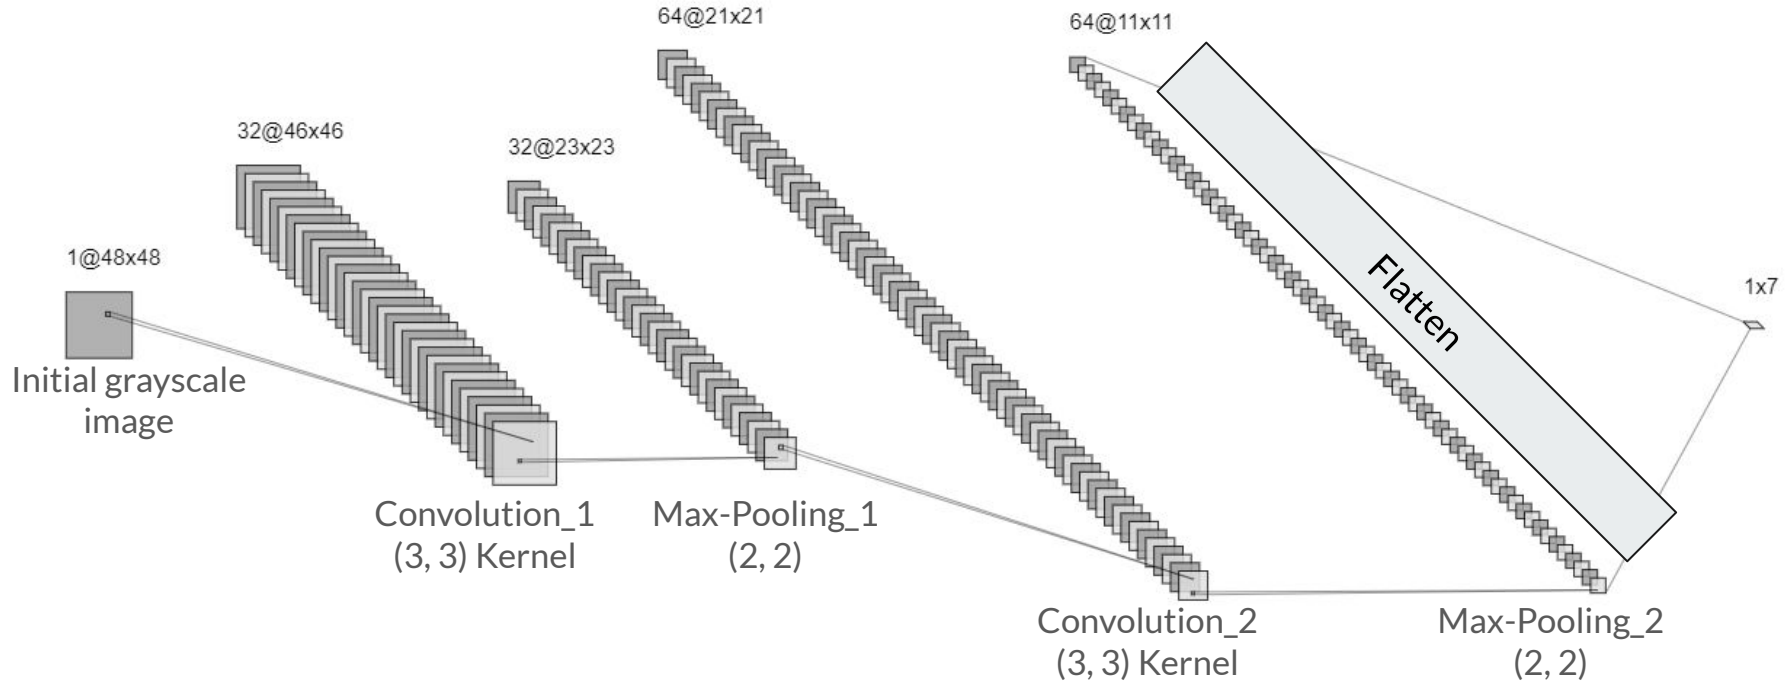
- Human-computer interaction (ex: face id)
- Applicable to mental health tools (detecting depression)
- Helpful with threat detection (security applications)

Both techniques are well documented approaches for facial recognition, and we will explore the strengths and weaknesses of both.

Distribution of Labels in Train and Test Sets



Convolutional Neural Network (CNN) V1



Data Preprocessing

```
transform = transforms.Compose([
    transforms.Grayscale(),
    transforms.Resize((48, 48)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.5], std=[0.5])
])

train_dataset = datasets.ImageFolder(root='data/train', transform=transform)
test_dataset = datasets.ImageFolder(root='data/test', transform=transform)

class_counts = {"angry": 3995, "disgust": 436, "fear": 4097, "happy": 7215, "neutral": 4965, "sad": 4830, "surprise": 3171}



# Sampled with weights inversely proportional to the number of images in each class
weights = [1.0 / counts for counts in class_counts.values()]
labels = train_dataset.classes
sample_weights = [weights[label] for _, label in train_dataset.samples]

# Create sampler and DataLoader for training
sampler = WeightedRandomSampler(sample_weights, len(sample_weights))

train_loader = DataLoader(dataset=train_dataset, batch_size=32, sampler=sampler, shuffle=False)
test_loader = DataLoader(dataset=test_dataset, batch_size=32, shuffle=False)
```

CNN V1


- Model does not learn past Epoch 2
- Underfitting? Overfitting?
- Validation Accuracy: 0.5181
- F1 Score: 0.5220



CNN V2



- Additional Convolutional Layers
- Batch Normalization
- Initially resulted in overfitting:
 - Dropout Layers
 - Increased dropout percentages
- Validation Accuracy: 0.6067
- F1 Score: 0.5887

```
FERModel(  
  (cnn_layers): Sequential(  
    (0): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): ReLU()  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Dropout(p=0.5, inplace=False)  
    (9): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (10): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (11): ReLU()  
    (12): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (13): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (15): ReLU()  
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (17): Dropout(p=0.5, inplace=False)  
  )  
  (output_layer): Linear(in_features=2304, out_features=7, bias=True)  
)
```

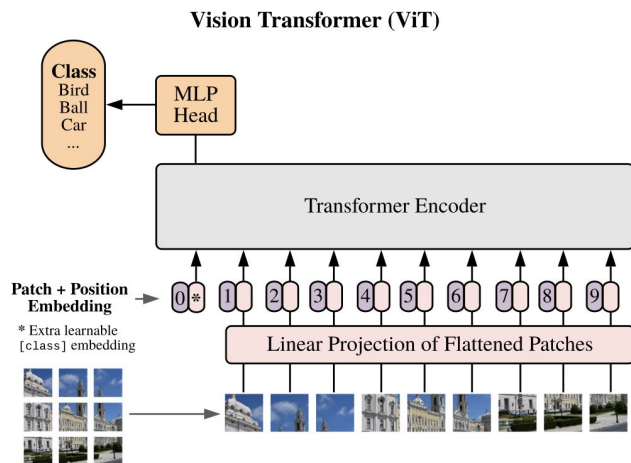


CNN V2.1 (Final)


- Same Architecture as V2
- Retrained on resampled



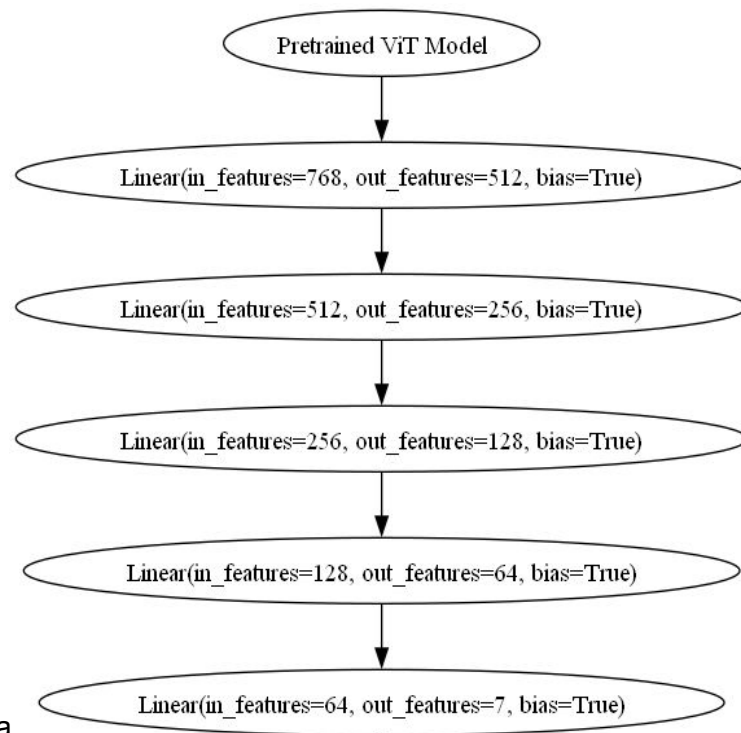
ViT Architecture



Transformer Encoder



Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, Neil Houlsby **“An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale”**



Initial attempt:


Train a ViT model from zero

Architecture: Convolutional Layer + Transformer Encoder + Fully Connected Layer

Result:


- Training loss didn't decrease (~1.8)
- Accuracy: 0.25

Reason: ViT doesn't fit small datasets



Using device: cuda

Epoch	Progress	Loss	Time	Batch/s
Epoch 1/20	100%	898/898	[01:34<00:00, 9.49batch/s]	
Epoch [1/20]		Loss: 1.8549, Time: 94.64 sec		
Epoch 2/20	100%	898/898	[00:20<00:00, 43.93batch/s]	
Epoch [2/20]		Loss: 1.8181, Time: 20.44 sec		
Epoch 3/20	100%	898/898	[00:20<00:00, 44.23batch/s]	
Epoch [3/20]		Loss: 1.8154, Time: 20.31 sec		
Epoch 4/20	100%	898/898	[00:20<00:00, 43.81batch/s]	
Epoch [4/20]		Loss: 1.8147, Time: 20.50 sec		
Epoch 5/20	100%	898/898	[00:20<00:00, 44.13batch/s]	
Epoch [5/20]		Loss: 1.8137, Time: 20.35 sec		
Epoch 6/20	100%	898/898	[00:20<00:00, 43.27batch/s]	
Epoch [6/20]		Loss: 1.8132, Time: 20.76 sec		
Epoch 7/20	100%	898/898	[00:22<00:00, 39.99batch/s]	
Epoch [7/20]		Loss: 1.8138, Time: 22.46 sec		
Epoch 8/20	100%	898/898	[00:21<00:00, 41.09batch/s]	
Epoch [8/20]		Loss: 1.8138, Time: 21.85 sec		
Epoch 9/20	100%	898/898	[00:24<00:00, 37.31batch/s]	
Epoch [9/20]		Loss: 1.8135, Time: 24.07 sec		
Epoch 10/20	100%	898/898	[00:22<00:00, 40.22batch/s]	
Epoch [10/20]		Loss: 1.8131, Time: 22.33 sec		
Epoch 11/20	100%	898/898	[00:22<00:00, 40.45batch/s]	
Epoch [11/20]		Loss: 1.8128, Time: 22.20 sec		
Epoch 12/20	100%	898/898	[00:22<00:00, 39.60batch/s]	
Epoch [12/20]		Loss: 1.8131, Time: 22.68 sec		
Epoch 13/20	100%	898/898	[02:20<00:00, 6.38batch/s]	
Epoch [13/20]		Loss: 1.8122, Time: 140.79 sec		
Epoch 14/20	100%	898/898	[00:25<00:00, 34.62batch/s]	
Epoch [14/20]		Loss: 1.8121, Time: 25.94 sec		
Epoch 15/20	100%	898/898	[00:19<00:00, 44.49batch/s]	



Using Pre-trained ViT Model



google/vit-base-patch16-224

Vision Transformer (ViT) model pre-trained on ImageNet-21k (14 million images, 21,843 classes) at resolution 224x224, and fine-tuned on ImageNet 2012 (1 million images, 1,000 classes) at resolution 224x224.

[google/vit-base-patch16-224](https://huggingface.co/google/vit-base-patch16-224) · Hugging Face

Data Preprocessing



The resolution of our test data images don't fit the ViT model, so we upscale the 48x48 images to 224x224. This may result in lower accuracies.


```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]),
])

train_dataset = datasets.ImageFolder(root='archive/train', transform=transform)
test_dataset = datasets.ImageFolder(root='archive/test', transform=transform)


train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
```

Learning Rate


 $1e-3$ doesn't converge



$1e-4$ just right



$1e-5$ converge too slowly



Fine-tuning - V1.3



Replace classifier layer in
pretrained model with
“CustomClassifier”

Accuracy: 0.59

10 epochs

```
class CustomClassifier(nn.Module):
    def __init__(self, input_dim, num_classes):
        super(CustomClassifier, self).__init__()
        self.fc1 = nn.Linear(input_dim, 512)
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, num_classes)
        self.dropout = nn.Dropout(0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc2(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc3(x)
        return x
```

Fine-tuning - V1.5



Accuracy: 0.61

Add 2 more fully connected layers

10 epochs

```
class CustomClassifier(nn.Module):
    def __init__(self, input_dim, num_classes):
        super(CustomClassifier, self).__init__()
        self.fc1 = nn.Linear(input_dim, 512)
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, 128)
        self.fc4 = nn.Linear(128, 64)
        self.fc5 = nn.Linear(64, num_classes)
        self.dropout = nn.Dropout(0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc2(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc3(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc4(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc5(x)
        return x
```

Fine-tuning - V1.6



Accuracy: 0.625

24 epochs


Gradually Unfreeze

- Don't update pretrained layers until achieving better accuracy


Notice overfitting on later epochs

Early stopping used

```
# Define a schedule to unfreeze layers gradually
unfreeze_schedule = {
    2: 'vit.encoder.layer.11',
    4: 'vit.encoder.layer.10',
    6: 'vit.encoder.layer.9',
    8: 'vit.encoder.layer.8',
    10: 'vit.encoder.layer.7',
    12: 'vit.encoder.layer.6',
    14: 'vit.encoder.layer.5',
    16: 'vit.encoder.layer.4',
    18: 'vit.encoder.layer.3',
}
```



Confusion Matrix



Fine-tuning - V1.6.6



Accuracy: 0.634

24 epochs

Used dataset upscale

Predicts “disgust” much better


```
train_dataset = datasets.ImageFolder(root='archive/train', transform=transform)
test_dataset = datasets.ImageFolder(root='archive/test', transform=transform)

class_counts = {"angry": 3995, "disgust": 436, "fear": 4097, "happy": 7215, "neutral": 4965, "sad": 4830, "surprise": 3171}


weights = [1.0 / counts for counts in class_counts.values()]
labels = train_dataset.classes
sample_weights = [weights[label] for _, label in train_dataset.samples]

# Create sampler and DataLoader for training
sampler = WeightedRandomSampler(sample_weights, len(sample_weights))

train_loader = DataLoader(dataset=train_dataset, batch_size=32, sampler=sampler, shuffle=False)
test_loader = DataLoader(dataset=test_dataset, batch_size=32, shuffle=False)
```




Confusion Matrix



Fine-tuning - V1.6.7

(Final version)

Train and Validation Loss vs. Epochs



Accuracy: 0.663

24 epochs

Added batch normalization layers

```
class CustomClassifier(nn.Module):
    def __init__(self, input_dim, num_classes):
        super(CustomClassifier, self).__init__()
        self.fc1 = nn.Linear(input_dim, 512)
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, 128)
        self.fc4 = nn.Linear(128, 64)
        self.fc5 = nn.Linear(64, num_classes)
        self.bn0 = nn.BatchNorm1d(input_dim)
        self.bn1 = nn.BatchNorm1d(512)
        self.bn2 = nn.BatchNorm1d(256)
        self.bn3 = nn.BatchNorm1d(128)
        self.bn4 = nn.BatchNorm1d(64)
        self.dropout = nn.Dropout(0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        #x = self.bn0(x)

        x = self.fc1(x)
        x = self.bn1(x)
        x = self.relu(x)
        x = self.dropout(x)


        x = self.fc2(x)
        x = self.bn2(x)
        x = self.relu(x)
        x = self.dropout(x)

        x = self.fc3(x)
        x = self.bn3(x)
        x = self.relu(x)
        x = self.dropout(x)

        x = self.fc4(x)
        x = self.bn4(x)
        x = self.relu(x)
        x = self.dropout(x)

        x = self.fc5(x)
        return x
```

Confusion Matrix




Generalization



Input Image



Predicted Probabilities by Category






Input Image



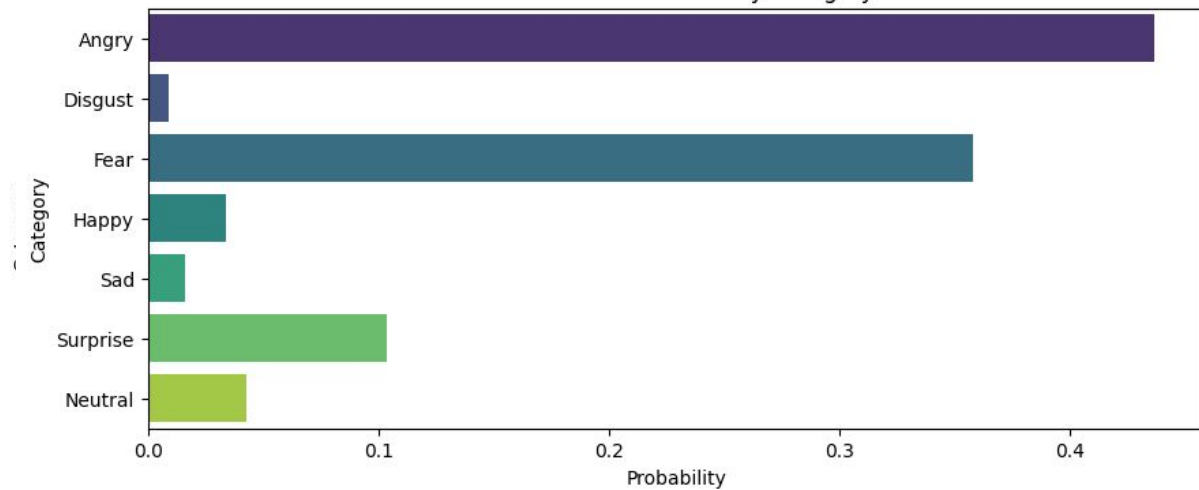
Predicted Probabilities by Category



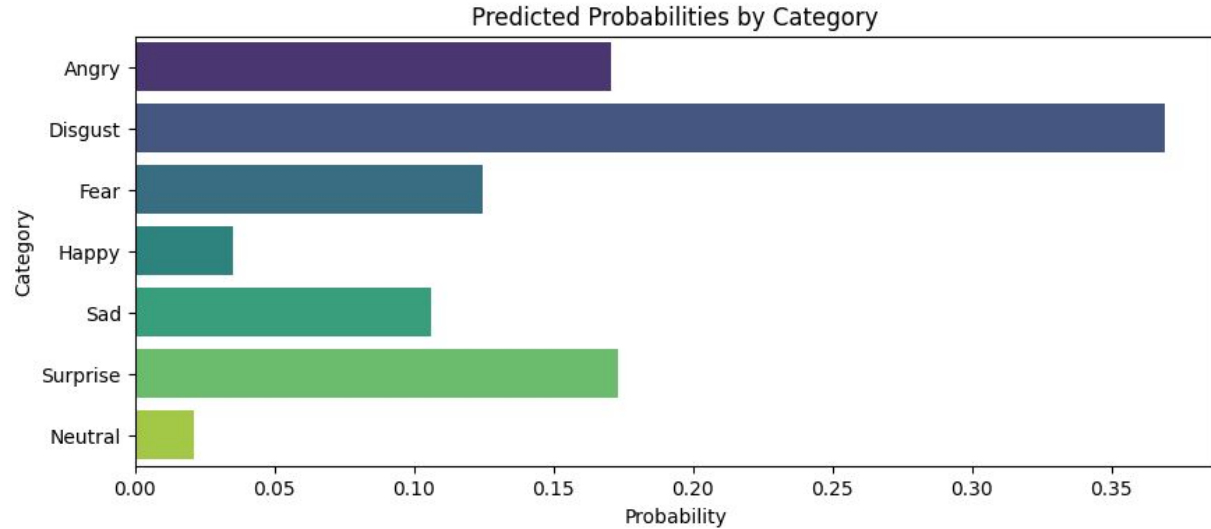
Input Image



Predicted Probabilities by Category




Input Image



One potential application of the model:

Understand your partner :-)

Interesting Discovery



“Loss can be seen as a **distance** between the true values of the problem and the values predicted by the model. The larger the loss, the larger the errors you made on the data.

Accuracy can be seen as the **count** of mistakes/ misclassifications you made on the data. The larger the accuracy, the fewer misclassifications you made on the data.”


<https://datascience.stackexchange.com/questions/42599/what-is-the-relationship-between-the-accuracy-and-the-loss-in-deep-learning>

Conclusion

We can see that the ViT performs better in 5 out of the 7 categories

However, the CNN does better in categories with fewer training samples (e.g. disgust)

Overall, it may be better to use a CNN when there is a small amount of training data/small image resolution, and a ViT otherwise



Note: The ViT is pre-trained on color images, while the test dataset is grayscale, this might be the cause of lower accuracies



Conclusion Continued

Potential problems with results:

- Google's ViT is pretrained and fine tuned on more data than our CNN architecture
 - unfair conclusion to say Vision Transformers are better than CNNs for this task
- The ViT might be performing worse since it was pretrained on higher resolution, colored images

Note: The ViT is pre-trained on color images, while the test dataset is grayscale, this might be the cause of lower accuracies



Future Works for ViT Model

How to combat overfitting?

- Freeze the pre-trained layers
- Use L2 regularization (weight decay)
- Adjusting the dropping out rate
- Use other pretrained models



Future Works

- Compare pre-trained CNN models to our ViT for a more fair comparison
- Find colored, upscaled image dataset for emotion recognition and rerun experiments



Q&A